

The Choice of Normalization Influences Shrinkage in Regularized Regression

Johan Larsson^{1,2,*} Jonas Wallin²

¹Department of Mathematical Sciences, University of Copenhagen, Denmark; ²Department of Statistics, Lund University, Sweden

Published in *Transactions on Machine Learning Research* (2025), Featured Certification



Take-Home Message

Normalization is **not** neutral preprocessing. For **binary features**, the *class balance* (proportion of ones) influences your lasso, ridge, and elastic net coefficients—and *which* normalization you pick decides *how* it does so. **The good news:** the right choice is predictable. Match the scaling (δ) to your penalty and goal, or tune it.

The Problem

Regularized models are scale-sensitive, so we *normalize* (center and scale) features first. But there are many ways to do this, and for binary features there is no obvious choice. Despite this, the effect of that choice has gone unstudied.

Figure 1 shows lasso paths for a real data set (w1a) under two types of normalization. The normalization choice alone decides which features look important.




Figure 1: Lasso paths for w1a under standardization ($\delta = 0.5$) and max-abs ($\delta = 0$), with first features to enter highlighted.

Setup

We study the elastic net (Zou & Hastie, 2005),

$$\frac{1}{2} \|\mathbf{y} - \beta_0 - \tilde{\mathbf{x}}\boldsymbol{\beta}\|_2^2 + \lambda_1 \|\boldsymbol{\beta}\|_1 + \frac{\lambda_2}{2} \|\boldsymbol{\beta}\|_2^2,$$

which gives the **lasso** ($\lambda_2 = 0$) and **ridge** ($\lambda_1 = 0$) as special cases. Each binary feature $\mathbf{x}_j$ is scaled by $s_j = (q_j - q_j^2)^\delta$, $\delta \geq 0$, covering no scaling ($\delta = 0$), standard deviation ($\delta = 0.5$), and variance ($\delta = 1$).

Class Balance and Estimation

With orthogonal features and no noise, the estimate on the original scale is

$$\hat{\beta}_j = \frac{\text{ST}_{\lambda_1} \left(\frac{\beta_j^* n (q_j - q_j^2)}{s_j} \right)}{s_j \left(\frac{n (q_j - q_j^2)}{s_j^2} + \lambda_2 \right)}.$$

The feature variance $q_j - q_j^2$ couples class balance to the estimate, mediated by the scale s_j . The right δ makes it cancel (Table 1).

Table 1: The δ that cancels class-balance bias for each penalty.

Penalty	Canceling Scale	δ
Lasso	Variance	1
Ridge	Standard deviation	$\frac{1}{2}$
Elastic net	None exists!	

For the elastic net, *no* choice removes the bias. Instead weight the *penalty*.

Selection Follows δ

Lower δ means imbalanced features are far less likely to be selected. As $q \rightarrow 1$, selection probability $\rightarrow 0$ for $\delta < \frac{1}{2}$ and $\rightarrow 1$ for $\delta > \frac{1}{2}$; see Figure 2.




Figure 2: Lasso selection probability vs. class balance q for several δ . Dotted line: asymptotic limit at $\delta = \frac{1}{2}$.

Beyond Orthogonality

The class-balance bias is not just an artifact of the orthogonal theory. It persists with many noisy, correlated features.

With 20 signals whose class balance rises $0.5 \rightarrow 0.99$, no scaling ($\delta = 0$) shrinks the most imbalanced signals to zero, while variance scaling ($\delta = 1$) recovers them (Figure 3).




Figure 3: Lasso estimates (mean $\pm$ SD) for a simulation with $p = 1000$ (the first 20 features are signals with class balance $0.5 \rightarrow 0.99$; first 30 shown). No scaling drops the imbalanced signals; variance scaling recovers them.

Tuning δ

Larger δ removes class-balance bias but inflates variance, so there is no universal best δ . On real data, the cross-validated error is minimized at an *interior* δ that differs by data set—so treat δ as a hyperparameter and tune it (Figure 4).




Figure 4: Hold-out (validation) NMSE over (λ, δ) for the lasso on three data sets. White dot: tuned optimum; dotted line: best δ at each λ . The optimal δ is interior and data set-dependent.

Recommendations

Selection

If missing a feature is costly, scale by *variance* ($\delta = 1$).

Prediction

Standard deviation or no scaling; better, *tune* δ as a hyperparameter.

Elastic net

Weight the *penalty*, not the features.

Why It Matters

The two most common defaults are *both* biased: standardization (the glmnet default; Friedman et al., 2010) and no scaling (recommended for sparse data in scikit-learn). The class-balance effect is therefore likely pervasive in published analyses.

References

- Zou & Hastie (2005). Regularization and variable selection via the elastic net. *JRSS-B*.
- Gelman (2008). Scaling regression inputs by dividing by two standard deviations. *Statistics in Medicine*.
- Friedman, Hastie & Tibshirani (2010). Regularization paths for generalized linear models via coordinate descent. *JSS*.
- Bien, Taylor & Tibshirani (2013). A lasso for hierarchical interactions. *Annals of Statistics*.
- Lim & Hastie (2015). Learning interactions via hierarchical group-lasso regularization. *JCGS*.

Contact

Email

johan@jolars.co

Homepage


<https://jolars.co>

Mastodon

[@jolars@mastodon.social](https://jolars@mastodon.social)

Bluesky

[@jolars.co](https://jolars.co)



Read the paper!